

Instructor 2 - Seaborn 2022

June 20, 2022

```
[66]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline
# %matplotlib is a magic function in IPython.

df = sns.load_dataset("diamonds")
diamonds = sns.load_dataset("diamonds")
tips = sns.load_dataset("tips")
penguins = sns.load_dataset("penguins")
flights = sns.load_dataset("flights")

df.shape
```

```
[66]: (53940, 10)
```

```
[2]: # sns.get_dataset_names()
```

1 Matplotlib to Pandas to Seaborn


```
[3]: # NOTICE: One y value per X value

x = ['Ideal', 'Premium', 'Very Good', 'Good', 'Fair',]
y = [3457, 4584, 3981, 3928, 4358]
```

```
[4]: type(x)
```


```
[4]: list
```

```
[5]: # Matplotlib
plt.plot(x, y);
```



```
[64]: # This works but data is not aggregated  
plt.plot(df.cut, df.price)
```

```
[64]: [<matplotlib.lines.Line2D at 0x7f77c979fdc0>]
```



```
[65]: df2 = df.groupby('cut').mean().reset_index() # creating a smaller dataset
df2
```

```
[65]:
```

	cut	carat	depth	table	price	x	y \
0	Ideal	0.702837	61.709401	55.951668	3457.541970	5.507451	5.520080
1	Premium	0.891955	61.264673	58.746095	4584.257704	5.973887	5.944879
2	Very Good	0.806381	61.818275	57.956150	3981.759891	5.740696	5.770026
3	Good	0.849185	62.365879	58.694639	3928.864452	5.838785	5.850744
4	Fair	1.046137	64.041677	59.053789	4358.757764	6.246894	6.182652

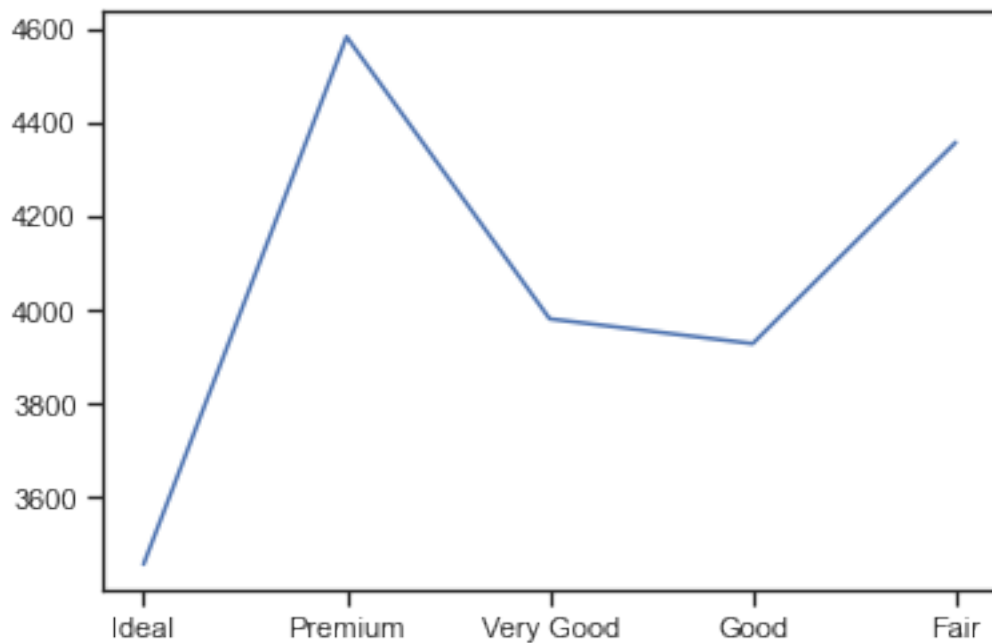

```

      z
0  3.401448
1  3.647124
2  3.559801
3  3.639507
4  3.982770

```

```
[69]: plt.plot(df2.cut, df2.price)
```


```
[69]: [<matplotlib.lines.Line2D at 0x7f782783e730>]
```



Below ... **pandas** plotting functions (e.g., `.plot`) uses **matplotlib**

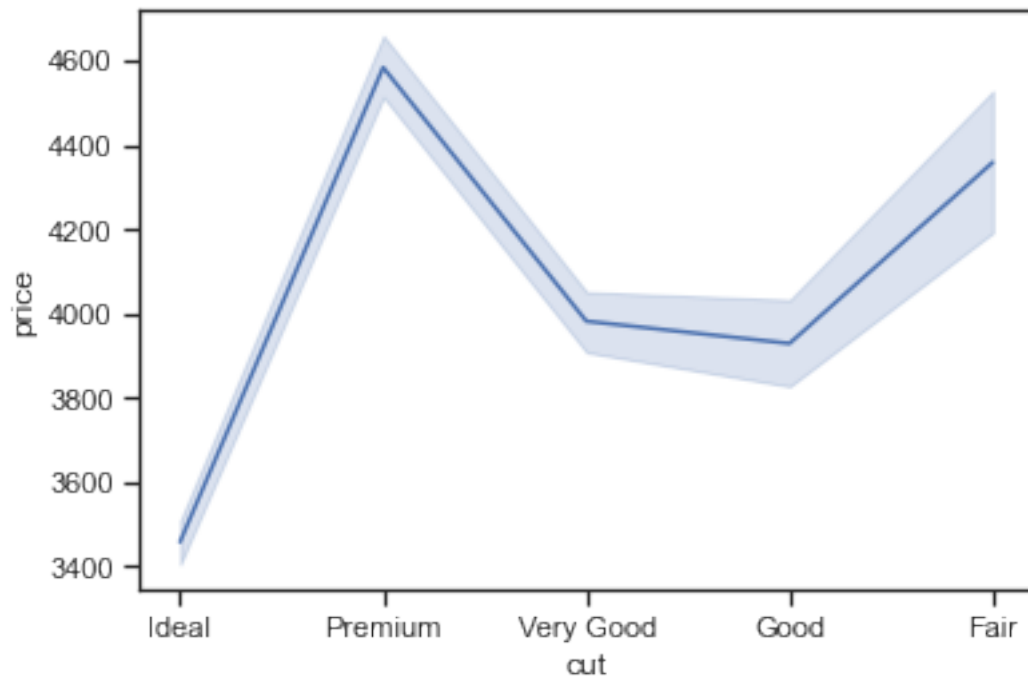
```
[72]: # This is a pandas plot of the raw data (df)
#df.plot(x='cut', y = 'price')
df2.plot(x='cut', y = 'price')
```

```
[72]: <AxesSubplot:xlabel='cut'>
```



```
[70]: # This is a seaborn plot of df (not grouped) - automatic aggregation
sns.lineplot(data=df, x="cut", y="price")
```

```
[70]: <AxesSubplot:xlabel='cut', ylabel='price'>
```



```
[10]: # Using seaborn's set_theme method
# Seaborn is updating Matplotlib's rc parameters. rc params are the default
# style settings
# rc = runtime configuration.
# https://matplotlib.org/stable/tutorials/introductory/customizing.html
```

```
sns.set_theme()
```

```
[11]: # Matplotlib
plt.plot(x, y);
```



2 Seaborn

Seaborn is a library for making statistical graphics in Python. It builds on top of matplotlib and integrates closely with pandas data structures.

Its plotting functions operate on dataframes and arrays containing whole datasets and internally perform the necessary semantic mapping and **statistical aggregation** to produce informative plots.

The **declarative** API lets you focus on what the different elements of your plots mean, rather than on the details of how to draw them.

2.1 Figure & Axes Level Plotting Functions


<http://seaborn.pydata.org/tutorial.html>

2.2 What's the difference between figure and axes level options?

Axes-level functions make self-contained plots The axes-level functions are written to act like drop-in replacements for matplotlib functions. While they add axis labels and legends automatically, they don't modify anything beyond the axes that they are drawn into. That means they can be composed into arbitrarily-complex matplotlib figures with predictable results.


```
[12]: # Figure level
sns.displot( data=df, x="price", hue="cut", multiple = "stack", kind = 'hist',
    ↪ bins = 20)
```

```
[12]: <seaborn.axisgrid.FacetGrid at 0x7f78263dc940>
```



```
[13]: # Axes level
sns.histplot(data=df, x='price', hue='cut', multiple = 'stack', bins = 20)
```




```
[13]: <AxesSubplot:xlabel='price', ylabel='Count'>
```



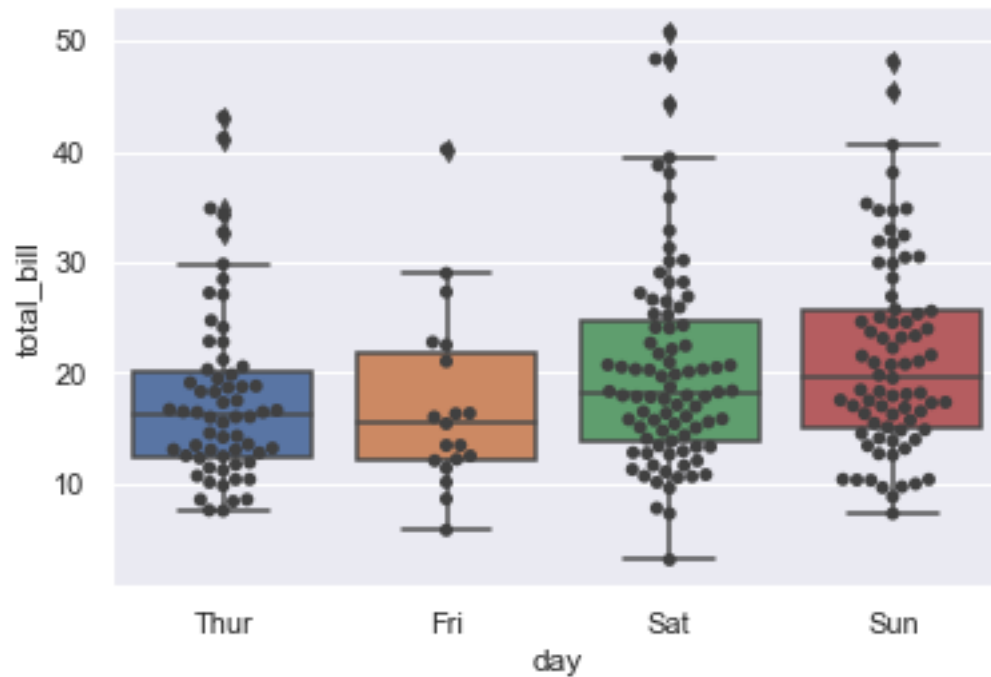
2.3 Using sns and matplotlib together

```
[14]: plt.figure(figsize=(20, 15))
plt.subplot(2,2,1)
sns.boxplot(y = 'depth', x = 'cut', data = df)
plt.subplot(2,2,2)
sns.scatterplot(y = 'price', x = 'carat', data = df)
plt.subplot(2,2,3)
sns.boxplot(y = 'carat', x = 'cut', data = df)
plt.subplot(2,2,4)
sns.histplot(x = 'price', data = df)
```

```
[14]: <AxesSubplot:xlabel='price', ylabel='Count'>
```

```
[15]: # From Seaborn documentation
ax = sns.boxplot(x="day", y="total_bill", data=tips)
ax = sns.swarmplot(x="day", y="total_bill", data=tips, color=".25")
```




2.4 Figure level

- Figure-level functions interface with matplotlib through a seaborn object, usually a FacetGrid
- Each module (relational, distributions, categorical) has a single figure-level function

```
[16]: # The default for distplot is a histogram


sns.displot(data=df, x="price")
plt.savefig('save_as_a_png.png')

# sns.displot(data=df, x="price", height=8, aspect=15/8) Use height and aspect_
↪ to change the size of the figure.
```



```
[17]: sns.displot(data=df, x="price", hue="cut", multiple="stack")
```


```
[17]: <seaborn.axisgrid.FacetGrid at 0x7f7827591be0>
```



2.5 Change the plot type with kind = ...

```
[18]: sns.displot(data=df, x="price", hue="cut", col="cut", kind = 'hist')
```


```
[18]: <seaborn.axisgrid.FacetGrid at 0x7f7809b6fe50>
```



```
[19]: sns.displot(data=df, x="price", hue="cut", col="cut", kind = 'kde')
```

```
# kernel density estimation
```


```
[19]: <seaborn.axisgrid.FacetGrid at 0x7f782640a400>
```



```
[20]: sns.displot(data=df, x="price", hue="cut", col="cut", kind = 'ecdf')


# empirical cumulative distribution functions
```

```
[20]: <seaborn.axisgrid.FacetGrid at 0x7f780a3dcf70>
```




```
[21]: sns.displot(data=df, x="price", hue="cut", kind = 'kde', rug = True)
```

```
[21]: <seaborn.axisgrid.FacetGrid at 0x7f780bdbca00>
```



```
[22]: # This one might take a minute to run.  
  
sns.displot(data=df, x="carat", y='price', kind='kde')
```

```
[22]: <seaborn.axisgrid.FacetGrid at 0x7f780c82f610>
```



2.6 Seaborn Exercise 1 - 10 minutes

- Use the relational (relplot) figure-level function to create two charts. First a scatterplot and second a line chart.
- Use the 'tips' data set.
- For the scatterplot, determine if tips increase with the bill amount. Try to show a distinction between data points based on time of day.
- For the line chart, show how tips change based on size of the party.


```
[23]: tips.head()
```

```
[23]:
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4


```
[24]: # Place scatterplot here
```

```
sns.relplot(x="total_bill", y="tip", data=tips, kind = 'scatter', hue = 'smoker');
```




```
[25]: # Place line chart here
```

```
sns.relplot(x="size", y="tip", data=tips, kind = 'line');
```





3 Facet Grids - Creating Small Multiples


```
[26]: p = sns.FacetGrid(df) # p is the facet grid
```




```
[27]: p = sns.FacetGrid(df, col = 'cut') # 1 column for each facet (value) of cut.
      # matplotlib will squeeze the 5 plots into the original size.
```



```
[28]: p = sns.FacetGrid(df, col = 'cut', height = 4, aspect = 0.75)
      # Aspect ratio of each facet, so that aspect * height gives the width of each
      ↪ facet.
```




```
[29]: p = sns.FacetGrid(df, col = 'cut', height = 4, aspect = 0.75, col_wrap = 3)
      # Aspect ratio of each facet, so that aspect * height gives the width of each
      ↪ facet.
```




```
[30]: sns.set_style('white')
      penguins = sns.load_dataset("penguins")
```

```
[31]: p = sns.FacetGrid(penguins, col='island');
```



```
[32]: p = sns.FacetGrid(penguins, row='island');
```



```
[33]: type(p)
```

```
[33]: seaborn.axisgrid.FacetGrid
```

3.1 Managing the Facet Grid


```
[34]: penguins.head()
```

```
[34]:  species      island  bill_length_mm  bill_depth_mm  flipper_length_mm  \
0  Adelie  Torgersen         39.1           18.7           181.0
1  Adelie  Torgersen         39.5           17.4           186.0
2  Adelie  Torgersen         40.3           18.0           195.0
3  Adelie  Torgersen          NaN           NaN            NaN
4  Adelie  Torgersen         36.7           19.3           193.0

   body_mass_g      sex
0      3750.0    Male
1      3800.0  Female
2      3250.0  Female
3          NaN     NaN
4      3450.0  Female
```

```
[35]: sns.displot(data=penguins, x="flipper_length_mm", col="island", kind = 'hist')
```

```
[35]: <seaborn.axisgrid.FacetGrid at 0x7f780bde5490>
```



3.2 Using methods of FacetGrid

These will be used when an axes level plot is used.


Three steps: - set up the FacetGrid - identify the plot type using `.map` or `.map_dataframe` - customize

3.2.1 .map()

1. Set up the facet grid (format the facets)
2. Describe what should be plotted in the grids
3. Add extras - labels, titles, etc.


To draw a plot on every facet, pass a function and the name of one or more columns in the dataframe to `FacetGrid.map()`

```
[36]: p = sns.FacetGrid(penguins, col='island')
p.map(sns.histplot, 'flipper_length_mm'); # Requires positional arguments, not
↳ named (x = 'flipper_length_mm')
```




3.2.2 .map_dataframe()

```
[37]: p = sns.FacetGrid(penguins, col='island')
p.map_dataframe(sns.histplot, x = 'flipper_length_mm'); # It is the x= that is
↳ different. This is the named argument
```



```
[38]: p = sns.FacetGrid(penguins, col='island', height = 6, aspect = 1)
p.map_dataframe(sns.scatterplot, y='bill_length_mm', x='bill_depth_mm');
```




[]:

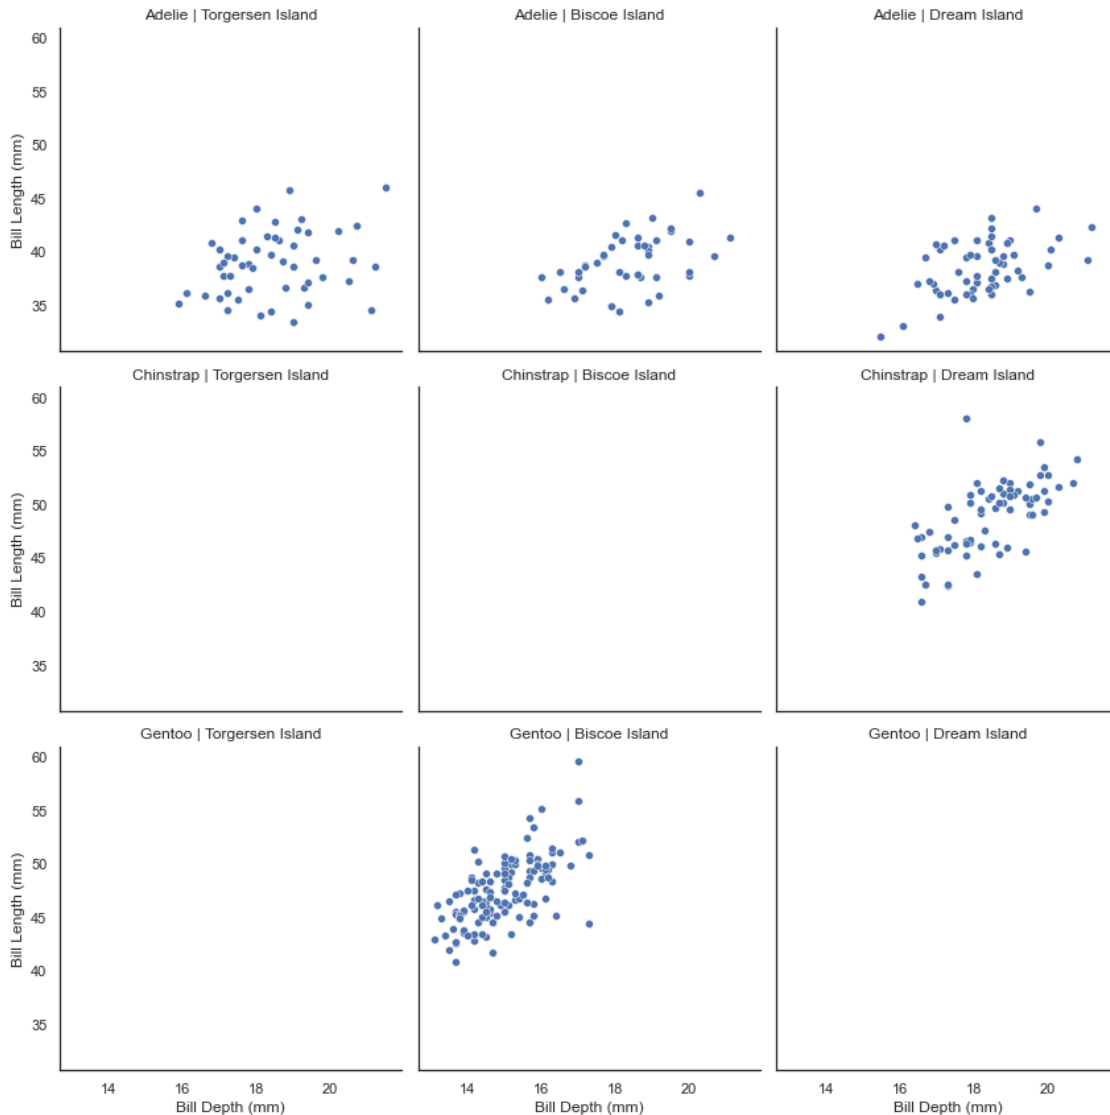
3.2.3 .set_axis_labels(), .set_titles(), sharey, ylim

```
[39]: p = sns.FacetGrid(penguins, col='island', height = 6, aspect = 1)
p.map_dataframe(sns.scatterplot, x='bill_depth_mm', y='bill_length_mm')

p.set_axis_labels('BILL DEPTH (mm)', 'BILL LENGTH (mm)'); # if the LABELS needs
↳ to be changed
p.set_titles(col_template='{col_name} Island'); # if the TITLE needs to be
↳ changed
```



```
[40]: p = sns.FacetGrid(penguins, col='island', row='species', height = 4, aspect = 1)
p.map_dataframe(sns.scatterplot, x='bill_depth_mm', y='bill_length_mm')
p.set_axis_labels('Bill Depth (mm)', 'Bill Length (mm)')
p.set_titles(row_template='{row_name}', col_template='{col_name} Island');
```


- sharey: False means the y-axis will not be shared and each plot will get its own y-axis.
- ylim: Sets a specified range for all y-axes shown

sharey = False

```
[41]: p = sns.FacetGrid(penguins, col='island', height = 4, aspect =1, sharey=False)


#p = sns.FacetGrid(penguins, col='island', height = 4, aspect =1, sharey=False,
↳ ylim=(20, 70))

p.map_dataframe(sns.scatterplot, x='bill_depth_mm', y='bill_length_mm');
p.set_axis_labels('BILL DEPTH (mm)', 'BILL LENGTH (mm)');
p.set_titles(col_template='{col_name} Island');
```




3.2.4 hue & palette

```
[42]: p = sns.FacetGrid(penguins, col='island', height = 4, aspect =1, sharey=False,
    ylim=(20, 70), hue = 'species')
p.map_dataframe(sns.scatterplot, x='bill_depth_mm', y='bill_length_mm');
p.set_axis_labels('BILL DEPTH (mm)', 'BILL LENGTH (mm)');
p.set_titles(col_template='{col_name} Island');
```



```
[43]: p = sns.FacetGrid(penguins,
    col='island',
    height = 4,
    aspect =1,
    sharey=False,
    ylim=(20, 70),
    hue = 'species',
    palette = 'magma'
    #palette = ['grey','blue','red']
)
```

```
p.map_dataframe(sns.scatterplot, x='bill_depth_mm', y='bill_length_mm', marker='+',
                 ↪=' ');
p.set_axis_labels('BILL DEPTH (mm)', 'BILL LENGTH (mm)');
p.set_titles(col_template='{col_name} Island');
```




3.2.5 Accomplish the same without defining the facet grid first

```
[44]: p = sns.relplot(data=penguins, x='bill_depth_mm', y='bill_length_mm', kind='scatter',
                    col='island', height = 4, aspect =1, hue = 'species', palette='magma',
                    marker = '+',
                    size = 'body_mass_g',
                    style = 'sex',
                    facet_kws={'sharey': False, 'sharex': True, 'ylim':(20,70)}
                    # sharey=False, ylim=(20, 70), #palette = ['grey','blue','red']
                    )

p.map(plt.axhline,
      y=45, color=".7",
      dashes=(2, 1),
      zorder=0)
p.set_axis_labels('Bill Depth (mm)', 'Bill Length (mm)')
p.set_titles(row_template='{row_name}', col_template='{col_name} Island')
```

```
[44]: <seaborn.axisgrid.FacetGrid at 0x7f7811063d60>
```




3.2.6 Method Chaining

```
[45]: p = sns.relplot(data=penguins, x='bill_depth_mm', y='bill_length_mm', kind = 'scatter',
    ↪ 'scatter',
    col='island', height = 4, aspect =1, hue = 'species', palette = 'magma'
    ↪ 'magma'
    ,marker = '+',
    size = 'body_mass_g',
    style = 'sex',
    facet_kws={'sharey': False, 'sharex': True, 'ylim':(20,70)})

(
p.map(plt.axhline,
    y=45, color=".7",
    dashes=(2, 1),
    zorder=0)
.set_axis_labels('Bill Depth (mm)', 'Bill Length (mm)')
.set_titles(row_template='{row_name}', col_template='{col_name} Island'))
```

[45]: <seaborn.axisgrid.FacetGrid at 0x7f78102cca60>



3.3 Seaborn Exercise 2 - 10 minutes

Using the flights info, create a visualization that plots - for each month - the number of passengers by year.

There should be one plot per month.

```
[46]: flights.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144 entries, 0 to 143
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   year        144 non-null    int64
1   month       144 non-null    category
2   passengers  144 non-null    int64
dtypes: category(1), int64(2)
memory usage: 2.9 KB
```

```
[47]: flights.head(20)
```

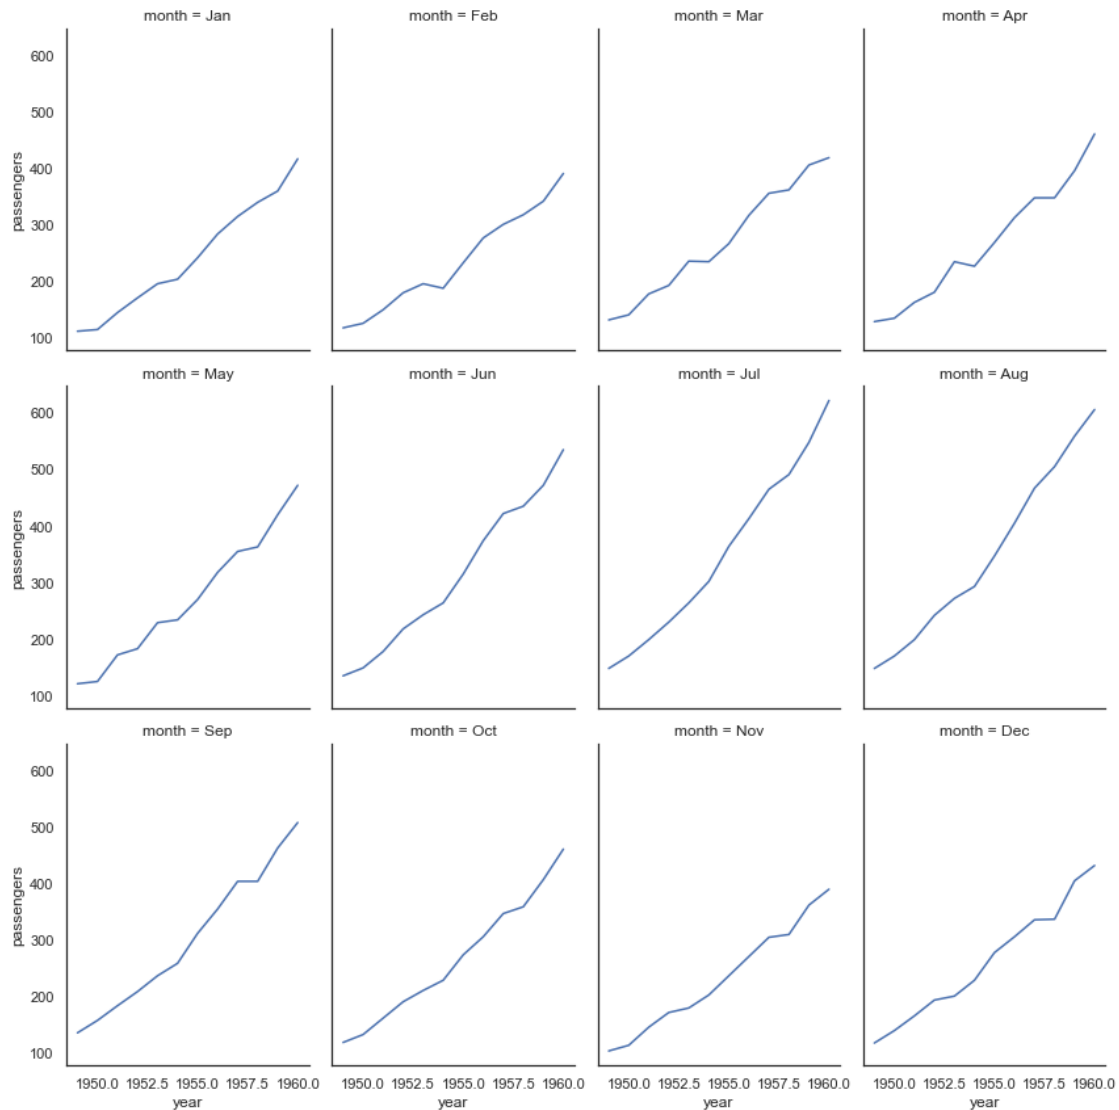
```
[47]:
```

	year	month	passengers
0	1949	Jan	112
1	1949	Feb	118
2	1949	Mar	132
3	1949	Apr	129
4	1949	May	121
5	1949	Jun	135
6	1949	Jul	148
7	1949	Aug	148
8	1949	Sep	136
9	1949	Oct	119
10	1949	Nov	104
11	1949	Dec	118
12	1950	Jan	115
13	1950	Feb	126
14	1950	Mar	141
15	1950	Apr	135
16	1950	May	125
17	1950	Jun	149
18	1950	Jul	170
19	1950	Aug	170

```
[48]: flights.shape
```

```
[48]: (144, 3)
```

```
[49]: # SNS Exercise 2 solution here.
p = sns.FacetGrid(flights, col = 'month', height = 4, aspect = 0.75, col_wrap = 4)
p.map_dataframe(sns.lineplot, x='year', y='passengers');
```




3.4 Seaborn Exercise 3 - 15 minutes

The distplot below is quick 'one-liner' plot. Take a little more time to create an axes for each cut and the axes are one above the other.

```
[50]: # This is the chart on the left with diamond data.
sns.displot(data=df, x="price", hue="cut", kind = 'kde', rug = True)
```

[50]: <seaborn.axisgrid.FacetGrid at 0x7f7812fb3b20>



```
[51]: # Place Exercise 3 solution here.

# https://towardsdatascience.com/
# ↪ sorry-but-sns-distplot-just-isnt-good-enough-this-is-though-ef2ddbf28078
df = sns.load_dataset("diamonds")

sns.set_style('white')
g = sns.FacetGrid(df, #the dataframe to pull from
                  row="cut", #define the column for each subplot row to be
                  ↪ differentiated by
                  hue="cut", #define the column for each subplot color to be
                  ↪ differentiated by
                  aspect=10, #aspect * height = width
                  height=1.5, #height of each subplot
                  palette=['#4285F4', '#EA4335', '#FBBC05', '#34A853'] #google
                  ↪ colors
                  )
```

```

#shade: True/False, shade area under curve or not
#alpha: transparency, lw: line width, bw: kernel shape specification

#g.map(sns.kdeplot, "price", lw=4, bw_method=0.2)  Same as below but no fill
g.map(sns.kdeplot, "price", shade=True, alpha=1, lw=1.5, bw_method=0.2)
g.map(plt.axhline, y=0, lw=4)


def label(x, color, label):
    ax = plt.gca() #get the axes of the current object
    ax.text(0, .2, #location of text
           label, #text label
           fontweight="bold", color=color, size=20, #text attributes
           ha="left", va="center", #alignment specifications
           transform=ax.transAxes) #specify axes of transformation

g.map(label, "x") #the function counts as a plotting object!

g.set_titles("") #set title to blank
g.set(yticks=[]) #set y ticks to blank
g.despine(bottom=True, left=True) #remove 'spines'

```

[51]: <seaborn.axisgrid.FacetGrid at 0x7f7825def7c0>




4 Multiple Views

4.1 Jointplot


```
[52]: sns.set_style("ticks")
sns.jointplot(data = penguins, x="flipper_length_mm", y="bill_length_mm",
             hue="species", height = 8 )
```

```
[52]: <seaborn.axisgrid.JointGrid at 0x7f7814eaf100>
```



```
[53]: sns.set_style("ticks")
sns.jointplot(data = penguins, x="flipper_length_mm", y="bill_length_mm",
             hue="species", height = 8 )
```

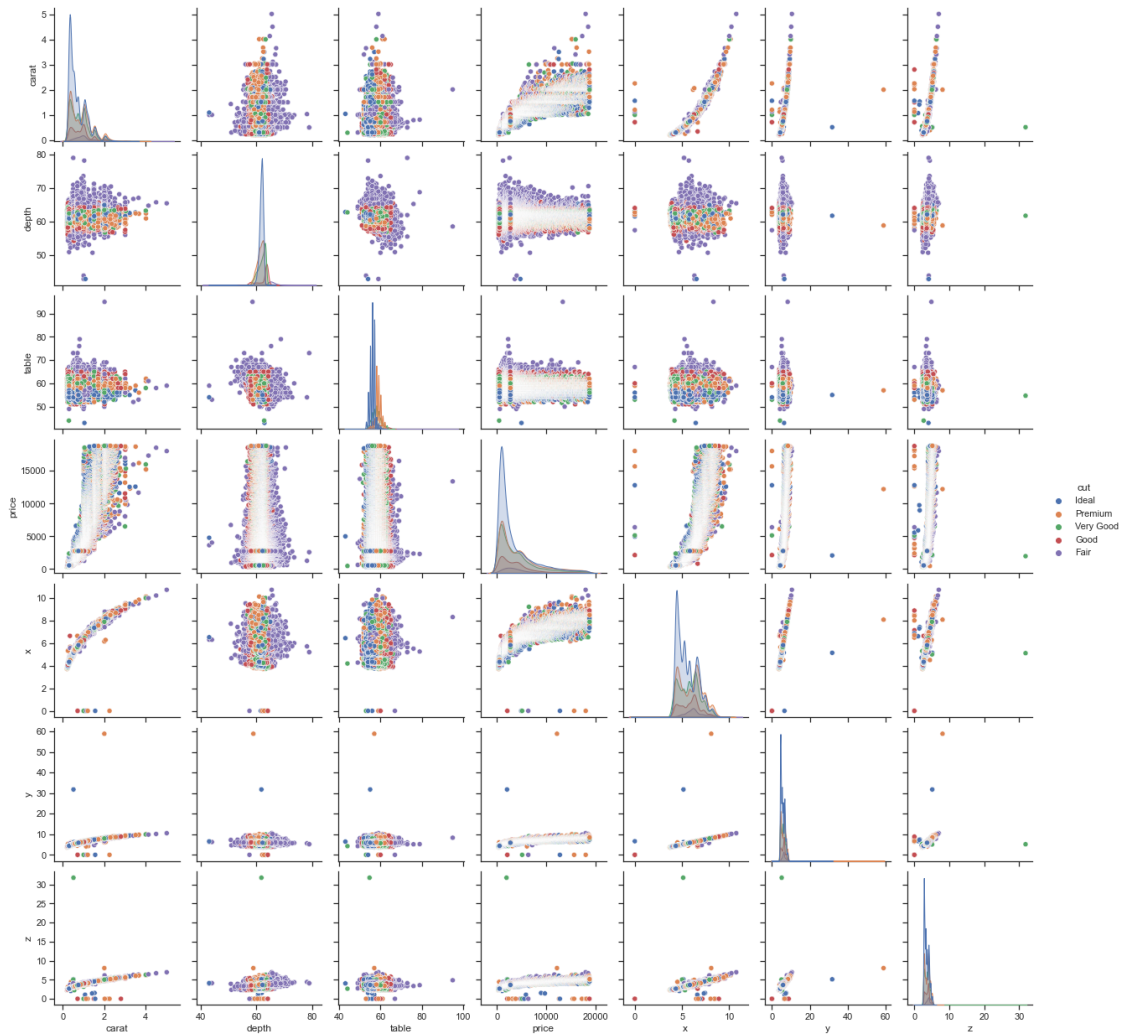
```
[53]: <seaborn.axisgrid.JointGrid at 0x7f7812258550>
```



4.2 Pairplot

```
[54]: sns.pairplot(data = df, hue = 'cut')
```

```
[54]: <seaborn.axisgrid.PairGrid at 0x7f7815107550>
```




```
[55]: xyz = df.corr()
      xyz
```

```
[55]:
```

	carat	depth	table	price	x	y	z
carat	1.000000	0.028224	0.181618	0.921591	0.975094	0.951722	0.953387
depth	0.028224	1.000000	-0.295779	-0.010647	-0.025289	-0.029341	0.094924
table	0.181618	-0.295779	1.000000	0.127134	0.195344	0.183760	0.150929
price	0.921591	-0.010647	0.127134	1.000000	0.884435	0.865421	0.861249
x	0.975094	-0.025289	0.195344	0.884435	1.000000	0.974701	0.970772
y	0.951722	-0.029341	0.183760	0.865421	0.974701	1.000000	0.952006
z	0.953387	0.094924	0.150929	0.861249	0.970772	0.952006	1.000000

```
[56]: sns.heatmap(xyz, annot=False)
```


```
[56]: <AxesSubplot:>
```




```
[57]: # Calculate correlations
corr = df.corr()
plt.figure(figsize=(12,8))
plt.title('Quantitative Variables Correlation')

# Heatmap
sns.heatmap(corr,cmap='plasma',annot=True)
```

```
[57]: <AxesSubplot:title={'center':'Quantitative Variables Correlation'}>
```



```
[58]: plt.figure(figsize=(12,8))
corr = df.corr()
ax = sns.heatmap(
    corr,
    vmin=-1, vmax=1, center=0,
    cmap=sns.diverging_palette(20, 220, n=200),
    square=True,
    annot=True, annot_kws={"size":10}
)
ax.set_xticklabels(
    ax.get_xticklabels(),
    rotation=45,
    horizontalalignment='right'
);
```



```
[59]: plt.figure(figsize=(12,8))
corr = df.corr()
ax = sns.heatmap(
    corr,
    vmin=-1, vmax=1, center=0,
    cmap=sns.diverging_palette(20, 220, n=200),
    square=True,
    annot=False, annot_kws={"size":20}
)
ax.set_xticklabels(
    ax.get_xticklabels(),
    rotation=45,
    horizontalalignment='right'
);
```

