

4 - Seaborn

March 22, 2023

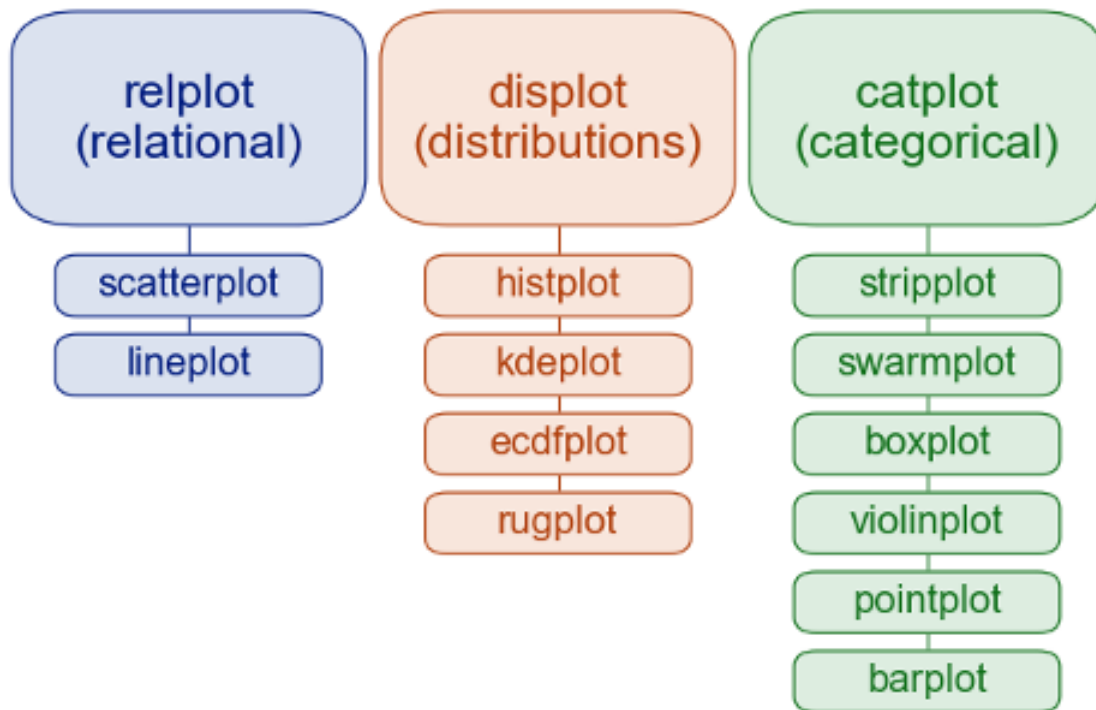
Table of Contents: Seaborn

- 1 Figure & Axes Level Plotting Functions
- 2 Figure level Functions
- 3 Axes level Functions
- 4 <http://seaborn.pydata.org/tutorial.html>
- 5 Figure level
- 6 Using the 'kind' kwarg
- 7 Seaborn Exercise 1 - 10 minutes
- 8 What's the difference?
 - 8.1 Combining matplotlib & seaborn syntax
- 9 Facet Grid
 - 9.1 `.map()`
 - 9.2 `.map_dataframe()`
 - 9.3 `.set_axis_labels()`, `.set_titles()`, `sharey`, `ylim`
 - 9.4 `hue` & `palette`
- 10 Multiple Views
 - 10.1 Jointplot
 - 10.2 Pairplot
- 11 Seaborn Exercise 2 - 10 minutes
- 12 Seaborn Exercise 3 - 15 minutes

Seaborn is a library for making statistical graphics in Python. It builds on top of matplotlib and integrates closely with pandas data structures.

Seaborn helps you explore and understand your data. Its plotting functions operate on dataframes and arrays containing whole datasets and internally perform the necessary semantic mapping and statistical aggregation to produce informative plots. Its dataset-oriented, declarative API lets you focus on what the different elements of your plots mean, rather than on the details of how to draw them.

0.1 Figure & Axes Level Plotting Functions



0.2 Figure level Functions

- `relplot`
- `displot` - default behavior is `histplot`
- `catplot`

0.3 Axes level Functions

- `scatterplot`
- `lineplot`
- `histplot`
- etc

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline

# import os
# for dirname, _, filenames in os.walk('/kaggle/input'):
#     for filename in filenames:
#         print(os.path.join(dirname, filename))
```

```
sns.set_theme()

#df = pd.read_csv('/kaggle/input/seaborn-practice/diamonds.csv')
#tips = pd.read_csv('/kaggle/input/seaborn-practice/tips.csv')
#penguins = pd.read_csv('/kaggle/input/seaborn-practice/penguins.csv')
#flights = pd.read_csv('/kaggle/input/seaborn-practice/flights.csv')

df = sns.load_dataset("diamonds")
tips = sns.load_dataset("tips")
penguins = sns.load_dataset("penguins")
flights = sns.load_dataset("flights")
```

```
[2]: # sns.get_dataset_names()
```

```
[3]: df.head()
```

```
[3]:
```

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43
1	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31
2	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31
3	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63
4	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75

```
[4]: tips.head()
```

```
[4]:
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

```
[5]: penguins.head()
```

```
[5]:
```

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	\
0	Adelie	Torgersen	39.1	18.7		181.0
1	Adelie	Torgersen	39.5	17.4		186.0
2	Adelie	Torgersen	40.3	18.0		195.0
3	Adelie	Torgersen	NaN	NaN		NaN
4	Adelie	Torgersen	36.7	19.3		193.0

	body_mass_g	sex
0	3750.0	Male
1	3800.0	Female
2	3250.0	Female
3	NaN	NaN
4	3450.0	Female

```
[6]: flights.head()
```


```
[6]:   year month  passengers
0  1949   Jan         112
1  1949   Feb         118
2  1949   Mar         132
3  1949   Apr         129
4  1949   May         121
```

0.4 <http://seaborn.pydata.org/tutorial.html>

0.5 Figure level


- Figure-level functions interface with matplotlib through a seaborn object, usually a FacetGrid
- Each module (relational, distributions, categorical) has a single figure-level function

```
[7]: # The default for distplot is a histogram
sns.distplot(data=df, x="price")
# plt.show() # removes the 'output' text
plt.savefig('save_as_a_png.png')
```



```
[8]: sns.displot(data=df, x="price", hue="cut", multiple="stack")
```


```
[8]: <seaborn.axisgrid.FacetGrid at 0x11333e490>
```



0.6 Using the 'kind' kwarg

```
[9]: sns.displot(data=df, x="price", hue="cut", col="cut", kind = 'hist')
```


```
[9]: <seaborn.axisgrid.FacetGrid at 0x128c81d30>
```



```
[10]: sns.displot(data=df, x="price", hue="cut", col="cut", kind = 'kde')

# kernel density estimation
```


```
[10]: <seaborn.axisgrid.FacetGrid at 0x129387040>
```



```
[11]: sns.displot(data=df, x="price", hue="cut", col="cut", kind = 'ecdf')


# empirical cumulative distribution functions
```

```
[11]: <seaborn.axisgrid.FacetGrid at 0x128f3a550>
```




```
[12]: sns.displot(data=df, x="price", hue="cut", kind = 'kde', rug = True)
```

```
[12]: <seaborn.axisgrid.FacetGrid at 0x129cb8d90>
```



```
[13]: # This one might take a minute to run.  
  
sns.displot(data=df, x="carat", y='price', kind='kde')
```

```
[13]: <seaborn.axisgrid.FacetGrid at 0x129ccf1f0>
```



0.7 Seaborn Exercise 1 - 10 minutes

- Use the relational (relplot) figure-level function to create two charts. First a scatterplot and second a line chart.
- Use the 'tips' data set.
- For the scatterplot, determine if tips increase with the bill amount. Try to show a distinction between data points based on time of day.
- For the line chart, show how tips change based on size of the party.

```
[14]: tips.head()
```

```
[14]:
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4


```
[15]: tips.shape
```

```
[15]: (244, 7)
```

```
[16]: # Place scatterplot here
```

```
[ ]:
```


```
[17]: # Place line chart here
```

```
#line aggregates data to the mean. If each data point is required, use the ↵  
↵index  
#sns.relplot(x=tips.index, y="tip", data=tips, kind = 'line', marker = 'h');
```

0.8 What's the difference?


```
[18]: sns.displot(data=df, x="price", hue="cut", multiple="stack", kind = 'hist', ↵  
↵bins = 20)
```

```
[18]: <seaborn.axisgrid.FacetGrid at 0x129cc0af0>
```



```
[19]: sns.histplot(data=df, x='price', hue='cut', multiple = 'stack', bins = 20)
```

```
[19]: <AxesSubplot:xlabel='price', ylabel='Count'>
```



Axes-level functions make self-contained plots The axes-level functions are written to act like drop-in replacements for matplotlib functions. While they add axis labels and legends automatically, they don't modify anything beyond the axes that they are drawn into. That means they can be composed into arbitrarily-complex matplotlib figures with predictable results.

0.8.1 Combining matplotlib & seaborn syntax

```
[20]: penguins.head()
```

```
[20]:
```

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	\
0	Adelie	Torgersen	39.1	18.7	181.0	
1	Adelie	Torgersen	39.5	17.4	186.0	
2	Adelie	Torgersen	40.3	18.0	195.0	
3	Adelie	Torgersen	NaN	NaN	NaN	
4	Adelie	Torgersen	36.7	19.3	193.0	

	body_mass_g	sex
0	3750.0	Male

```

1      3800.0  Female
2      3250.0  Female
3         NaN    NaN
4      3450.0  Female

```

```

[21]: # Example taken from Seaborn documentation

# Use penguins dataset


f, axs = plt.subplots(1, 2, figsize=(8, 4), gridspec_kw=dict(width_ratios=[4, 1,
↪3]))

sns.scatterplot(data=penguins,
                x="flipper_length_mm",
                y="bill_length_mm",
                hue="species",
                ax=axs[0])

sns.histplot(data=penguins,
             x="species",
             hue="species",
             shrink=.8,
             alpha=.8,
             legend=False,
             ax=axs[1])

f.tight_layout() # adjusts the space between subplots & around figure edge to
↪accomodate labels and content.


```



```
[22]: # Example taken from Seaborn documentation
# Use tips dataset


g = sns.relplot(data=tips, x="total_bill", y="tip")
g.ax.axline(xy1=(10, 2), slope=.2, color="b", dashes=(5, 2))
```

```
[22]: <matplotlib.lines._AxLine at 0x12bda6370>
```




0.9 Facet Grid


```
[23]: p = sns.FacetGrid(df)
```




```
[24]: p = sns.FacetGrid(df, col = 'cut')
      # matplotlib will squeeze the 5 plots into the original size.
```



```
[25]: p = sns.FacetGrid(df, col = 'cut', height = 4, aspect = 0.75)
      # Aspect ratio of each facet, so that aspect * height gives the width of each
      ↪ facet.
```




```
[26]: p = sns.FacetGrid(df, col = 'cut', height = 4, aspect = 0.75, col_wrap = 3)
      # Aspect ratio of each facet, so that aspect * height gives the width of each
      ↪ facet.
```



```
[27]: sns.set_style('white')
      penguins = sns.load_dataset("penguins")
```


```
[28]: p = sns.FacetGrid(penguins, col='island');
```



```
[29]: type(p)
```


```
[29]: seaborn.axisgrid.FacetGrid
```

```
[30]: p = sns.FacetGrid(penguins, row='island');
```




0.9.1 .map()

```
[31]: p = sns.FacetGrid(penguins, col='island')
p.map(sns.histplot, "flipper_length_mm");
```



0.9.2 .map_dataframe()

```
[32]: p = sns.FacetGrid(penguins, col='island')
p.map_dataframe(sns.histplot, y='flipper_length_mm');
```



```
[33]: penguins.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 344 entries, 0 to 343
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   species         344 non-null   object
```

```


1  island          344 non-null  object
2  bill_length_mm  342 non-null  float64
3  bill_depth_mm   342 non-null  float64
4  flipper_length_mm 342 non-null float64
5  body_mass_g     342 non-null  float64
6  sex            333 non-null  object
dtypes: float64(4), object(3)
memory usage: 18.9+ KB

```

```

[34]: p = sns.FacetGrid(penguins, col='island', height = 6, aspect =1)
p.map_dataframe(sns.scatterplot, x='bill_depth_mm', y='bill_length_mm');

```

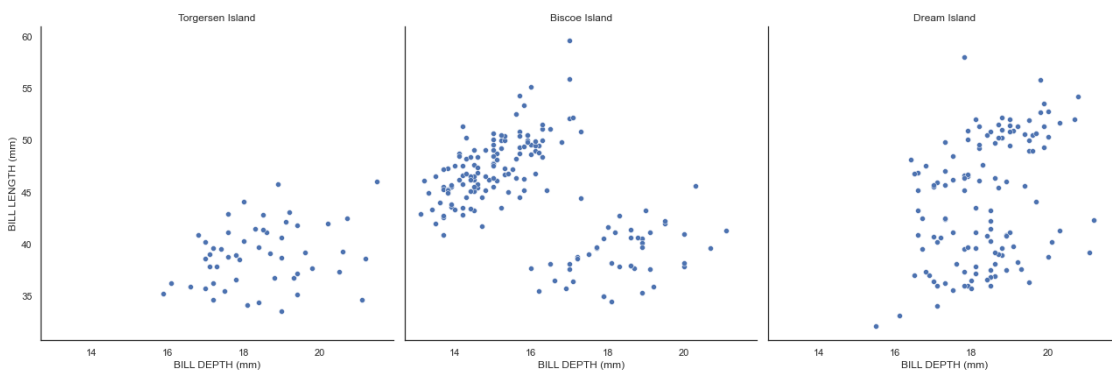


0.9.3 .set_axis_labels(), .set_titles(), sharey, ylim

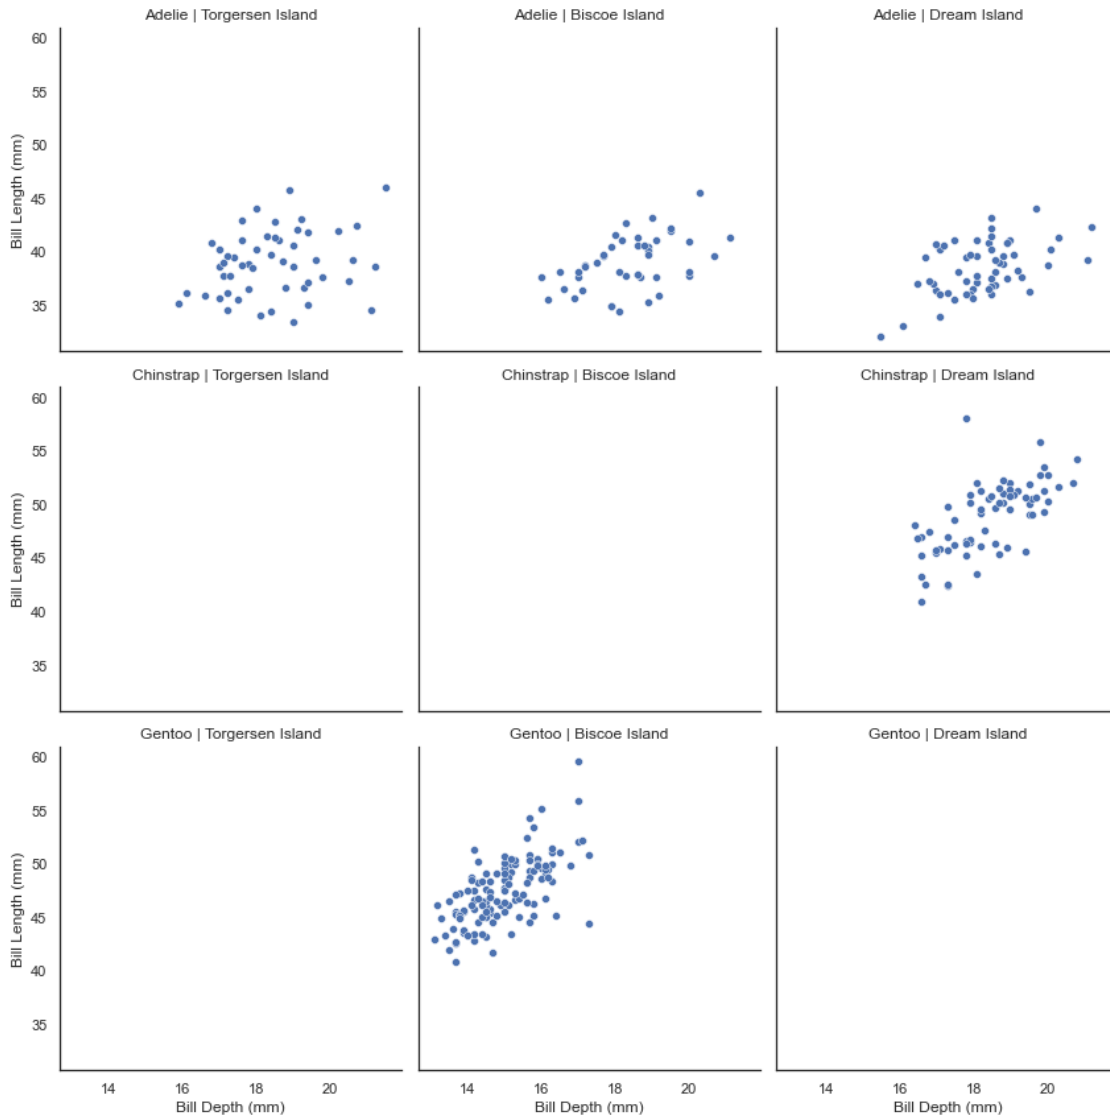
```

[35]: p = sns.FacetGrid(penguins, col='island', height = 6, aspect =1)
p.map_dataframe(sns.scatterplot, x='bill_depth_mm', y='bill_length_mm');
p.set_axis_labels('BILL DEPTH (mm)', 'BILL LENGTH (mm)'); # if the LABELS needs
↳to be changed
p.set_titles(col_template='{col_name} Island'); # if the TITLE needs to be
↳changed

```




```
[36]: p = sns.FacetGrid(penguins, col='island', row='species', height = 4, aspect =1)
p.map_dataframe(sns.scatterplot, x='bill_depth_mm', y='bill_length_mm')
p.set_axis_labels('Bill Depth (mm)', 'Bill Length (mm)')
p.set_titles(row_template='{row_name}', col_template='{col_name} Island');
```



- `sharey > False`: the y-axis will not be shared and each plot will get its own y-axis.
- `ylim >` Sets a specified range for all y-axes shown


```
[37]: p = sns.FacetGrid(penguins, col='island', height = 4, aspect =1, sharey=False)
#p = sns.FacetGrid(penguins, col='island', height = 4, aspect =1, sharey=False,
↳ ylim=(20, 70))
p.map_dataframe(sns.scatterplot, x='bill_depth_mm', y='bill_length_mm');
p.set_axis_labels('BILL DEPTH (mm)', 'BILL LENGTH (mm)');
```

```
p.set_titles(col_template='{col_name} Island');
```



0.9.4 hue & palette


```
[38]: p = sns.FacetGrid(penguins, col='island', height = 4, aspect =1, sharey=False,
    ylim=(20, 70), hue = 'species')
p.map_dataframe(sns.scatterplot, x='bill_depth_mm', y='bill_length_mm');
p.set_axis_labels('BILL DEPTH (mm)', 'BILL LENGTH (mm)');
p.set_titles(col_template='{col_name} Island');
```



Note: If hue is placed inside the scatterplot

```
[39]: p = sns.FacetGrid(penguins,
    col='island',
    height = 4,
    aspect =1,
    sharey=False,
    ylim=(20, 70),
    hue = 'species',
    palette = 'magma')
```

```
p.map_dataframe(sns.scatterplot, x='bill_depth_mm', y='bill_length_mm', marker='+',
                 hue='island');
p.set_axis_labels('BILL DEPTH (mm)', 'BILL LENGTH (mm)');
p.set_titles(col_template='{col_name} Island');
```




0.10 Multiple Views

0.10.1 Jointplot


```
[40]: sns.set_style("ticks")
sns.jointplot(data = penguins, x="flipper_length_mm", y="bill_length_mm",
              hue="species", height = 8 )
```

```
[40]: <seaborn.axisgrid.JointGrid at 0x12c71a0d0>
```



```
[41]: sns.set_style("ticks")
sns.jointplot(data = penguins, x="flipper_length_mm", y="bill_length_mm",
             hue="species", height = 8 )
```


```
[41]: <seaborn.axisgrid.JointGrid at 0x12c919070>
```



0.10.2 Pairplot

```
[42]: sns.pairplot(data = df, hue = 'cut')
```

```
[42]: <seaborn.axisgrid.PairGrid at 0x12ca33af0>
```




```
[43]: xyz = df.corr()
      xyz
```

```
[43]:
```

	carat	depth	table	price	x	y	z
carat	1.000000	0.028224	0.181618	0.921591	0.975094	0.951722	0.953387
depth	0.028224	1.000000	-0.295779	-0.010647	-0.025289	-0.029341	0.094924
table	0.181618	-0.295779	1.000000	0.127134	0.195344	0.183760	0.150929
price	0.921591	-0.010647	0.127134	1.000000	0.884435	0.865421	0.861249
x	0.975094	-0.025289	0.195344	0.884435	1.000000	0.974701	0.970772
y	0.951722	-0.029341	0.183760	0.865421	0.974701	1.000000	0.952006
z	0.953387	0.094924	0.150929	0.861249	0.970772	0.952006	1.000000

```
[44]: sns.heatmap(xyz, annot=False)
```


```
[44]: <AxesSubplot:>
```


```
[45]: # Calculate correlations
corr = df.corr()
plt.figure(figsize=(12,8))
plt.title('Quantitative Variables Correlation')

# Heatmap
sns.heatmap(corr,cmap='plasma',annot=True)
```


```
[45]: <AxesSubplot:title={'center':'Quantitative Variables Correlation'}>
```



```
[46]: plt.figure(figsize=(12,8))
corr = df.corr()
ax = sns.heatmap(
    corr,
    vmin=-1, vmax=1, center=0,
    cmap=sns.diverging_palette(20, 220, n=200),
    square=True,
    annot=True, annot_kws={"size":10}
)
ax.set_xticklabels(
    ax.get_xticklabels(),
    rotation=45,
    horizontalalignment='right'
);
```



```
[47]: plt.figure(figsize=(12,8))
corr = df.corr()
ax = sns.heatmap(
    corr,
    vmin=-1, vmax=1, center=0,
    cmap=sns.diverging_palette(20, 220, n=200),
    square=True,
    annot=False, annot_kws={"size":20}
)
ax.set_xticklabels(
    ax.get_xticklabels(),
    rotation=45,
    horizontalalignment='right'
);
```



0.11 Seaborn Exercise 2 - 10 minutes

Using the flights info, create a visualization that plots - for each month - the number of passengers by year.

There should be one plot per month.

```
[48]: flights.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144 entries, 0 to 143
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   year        144 non-null   int64
1   month       144 non-null   category
2   passengers  144 non-null   int64
dtypes: category(1), int64(2)
```

memory usage: 2.9 KB

```
[49]: flights.head(20)
```

```
[49]:
```

	year	month	passengers
0	1949	Jan	112
1	1949	Feb	118
2	1949	Mar	132
3	1949	Apr	129
4	1949	May	121
5	1949	Jun	135
6	1949	Jul	148
7	1949	Aug	148
8	1949	Sep	136
9	1949	Oct	119
10	1949	Nov	104
11	1949	Dec	118
12	1950	Jan	115
13	1950	Feb	126
14	1950	Mar	141
15	1950	Apr	135
16	1950	May	125
17	1950	Jun	149
18	1950	Jul	170
19	1950	Aug	170


```
[50]: flights.shape
```

```
[50]: (144, 3)
```

```
[51]: # Place solution here
```


0.12 Seaborn Exercise 3 - 15 minutes

The distplot below is quick ‘one-liner’ plot. Take a little more time to create an axes for each cut and the axes are one above the other.



```
[52]: sns.displot(data=df, x="price", hue="cut", kind = 'kde', rug = True)
```

```
[52]: <seaborn.axisgrid.FacetGrid at 0x12db30790>
```



```
[53]: # Place Exercise 3 solution here.
```

```
# How do we get 5 separate plots? How do we get each on a row?
```

```
# https://towardsdatascience.com/
```

```
→ sorry-but-sns-distplot-just-isnt-good-enough-this-is-though-ef2ddbf28078
```